

Introducing R

CJ 702: Advanced Criminal Justice Statistics

Thomas Bryan Smith*

February 03, 2025

Contents

1	Working Directories and Libraries	1
2	Objects	2
3	Vectors and Matrices	4
4	Logical Operations	9
5	Performing Simple Statistical Functions	14
6	Creating Custom Functions	17

1 Working Directories and Libraries

```
# Find the current working directory (cwd):  
getwd()
```

```
## [1] "c:/Users/Tom Smith/Documents/GitHub/CJStatsOM/1 Introducing R"
```

```
# ===== #  
  
# If the current working directory does not match the location of the folder  
# within which you want to work, you can change the directory with the  
# following code:  
  
# setwd("C:/the/working/directory")  
# Windows style formatting (note the forward slash)  
  
# setwd("/the/working/directory")
```

*University of Mississippi, tbsmit10@olemiss.edu

```

# Apple style formatting

# To run this code, you will need to replace the text between the
# speech marks, "", with your own working directory. This will
# depend on where you have saved your scripts, data, etc.

# The leading hashtag, #, "comments out" the code so that it is not
# interpreted by R. If you remove it, the code becomes "active"
# and will be interpreted.

# ===== #

# Having set your working directory, next you will want to install any packages
# you want to work with in your project(s). You do so with the following code:
# install.packages("tidyverse")

# Once the package is installed, you need to load it. You will need to do
# this in EVERY session, so the best practice is to include it at the
# beginning of each script.
library(tidyverse)

```

```
## Warning: package 'tidyverse' was built under R version 4.3.3
```

```
## Warning: package 'ggplot2' was built under R version 4.3.3
```

```
## Warning: package 'tidyr' was built under R version 4.3.3
```

```
## Warning: package 'readr' was built under R version 4.3.3
```

```
## Warning: package 'dplyr' was built under R version 4.3.3
```

```
## Warning: package 'stringr' was built under R version 4.3.3
```

```
## Warning: package 'lubridate' was built under R version 4.3.3
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
```

```
## v dplyr      1.1.4      v readr      2.1.5
```

```
## v forcats   1.0.0      v stringr   1.5.1
```

```
## v ggplot2    3.5.1      v tibble    3.2.1
```

```
## v lubridate  1.9.4      v tidyr     1.3.1
```

```
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

2 Objects

```
# R is an inherently mathematical language, and operates in a way that resembles  
# algebra. To "store" anything within R (text, single numbers, data, etc.)  
# you need to "assign" it to an "object". Once assigned, this object represents  
# the assigned value or data until it is overwritten, deleted, or you close the  
# session.
```

```
# Assign the numeric value, 8, to an object by the name "x":  
x <- 8
```

```
# Print the "x" object:  
print(x)
```

```
## [1] 8
```

```
# This can also be achieved by simply entering the object:  
x
```

```
## [1] 8
```

```
# You can also enter and print the assigned value simultaneously:  
(name <- "Mark")
```

```
## [1] "Mark"
```

```
# Depending on the type of the object,  
# you can perform mathematical operations:  
x + 2
```

```
## [1] 10
```

```
# These output of these mathematical operations can be assigned:  
(result <- x + 2)
```

```
## [1] 10
```

```
# Then you can continue to work with the new object:  
(result / 2)
```

```
## [1] 5
```

```
(result * 5)
```

```
## [1] 50
```

```
# You can save objects with the save() function:  
save(x, result, file = "myobjects.rda")  
# Note that this will save them to your current working directory!  
  
# You can view what is currently stored in your  
# "workspace" (all of your objects) with the ls() function:  
ls()
```

```
## [1] "name" "result" "x"
```

```
# Then the rm() function can be used to remove objects from that workspace:  
rm(x, result, name)
```

```
# You can also completely wipe your workspace of all objects:  
rm(list = ls())
```

```
# Now your workspace will be empty:  
ls()
```

```
## character(0)
```

```
# ...and you will not be able to "call" any of these objects:  
# result
```

```
# ...but you can load in any saved objects with the load() function:  
load(file = "myobjects.rda")
```

```
# The saved objects are now back in the workspace:  
ls()
```

```
## [1] "result" "x"
```

```
# Before moving on, let's wipe the workspace again:  
rm(list = ls())
```

3 Vectors and Matrices

```
# To manually creating a numeric vector, we use the  
# 'combine', c(), function:  
(x <- c(1, 2, 3, 4))
```

```
## [1] 1 2 3 4
```

```
# In this case, we can do the same vector by generating a sequence:  
y <- 1:4
```

```
# or, using the seq() function:  
z <- seq(from = 1, to = 4)
```

```
# The seq() function can also be used to generated more complex numeric vectors:  
seq(from = 1, to = 10, by = 2)
```

```
## [1] 1 3 5 7 9
```

```
# Then the rep() function can be used to generate vectors with repeating values:  
rep(2, times = 10)
```

```
## [1] 2 2 2 2 2 2 2 2 2 2
```

```
rep(1:2, times = 10)
```

```
## [1] 1 2 1 2 1 2 1 2 1 2 1 2 1 2 1 2 1 2
```

```
rep(1:2, length.out = 8)
```

```
## [1] 1 2 1 2 1 2 1 2
```

```
# Sometimes it can be helpful to create a sequence, starting at 1,  
# that iterates along another vector that does not start at 1:  
(x <- 51:65)
```

```
## [1] 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65
```

```
seq_along(x)
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15
```

```
# Vectors are not always numeric, they can be character vectors:  
(x <- c("Benjamin", "Kathryn", "Jean", "Jonathan"))
```

```
## [1] "Benjamin" "Kathryn" "Jean" "Jonathan"
```

```
# ...or they can be boolean (TRUE, FALSE) vectors:  
(x <- c(TRUE, FALSE, FALSE, TRUE))
```

```
## [1] TRUE FALSE FALSE TRUE
```

```
# ===== #  
# Mathematical operations function differently for vectors of length > 1.  
# Here are some examples:  
(x <- 1:4)
```

```
## [1] 1 2 3 4
```

```
(y <- 0:3)
```

```
## [1] 0 1 2 3
```

```
# [1 2 3 4] + 1  
1:4 + 1
```

```
## [1] 2 3 4 5
```

```
x + 1
```

```
## [1] 2 3 4 5
```

```
# [1 2 3 4] + [0 1]  
1:4 + 0:1
```

```
## [1] 1 3 3 5
```

```
x + 0.1
```

```
## [1] 1.1 2.1 3.1 4.1
```

```
# [1 2 3 4] + [0 1 2 3]  
1:4 + 0:3
```

```
## [1] 1 3 5 7
```

```
x + y
```

```
## [1] 1 3 5 7
```

```
# These rules also apply to other mathematical operations:  
x - y
```

```
## [1] 1 1 1 1
```

```
x * y
```

```
## [1] 0 2 6 12
```

```
x / y
```

```
## [1] Inf 2.000000 1.500000 1.333333
```

```
# Working with matrices is a little more complicated, and requires  
# knowledge of matrix algebra. Being as most of you do not work  
# with relational data, focusing more on single variables (vectors),  
# we will not discuss mathematical operations for matrices.
```

```
# ===== #
```

```
# Now that you understand vectors, let's move on to matrices.  
# Creating a simple matrix:  
(mat <- matrix(c(0,1,0, 15,18,10, 2,6,0, 126,75,50), nrow = 3))
```

```
##      [,1] [,2] [,3] [,4]  
## [1,]  0  15   2 126  
## [2,]  1  18   6  75  
## [3,]  0  10   0  50
```

```
# Note that the matrix is populated column-by-column. The nrow
# option defines the number of rows in the matrix. By splitting
# the input vector into chunks of 3, we can easily recognize
# what each column will look like ahead of time.
# When you change the number of rows, it helps to change the input:
(mat <- matrix(c(0,1,0,15, 18,10,2,6, 0,126,75,50), nrow = 4))
```

```
##      [,1] [,2] [,3]
## [1,]    0  18    0
## [2,]    1  10  126
## [3,]    0   2   75
## [4,]   15   6   50
```

```
# ...or, following best practice:
(mat <- matrix(c(0, 1, 0,      # column 1
                 15, 18, 10,   # column 2
                 2, 6, 0,      # column 3
                 126, 75, 50), # column 4
               nrow = 3))
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    0  15   2  126
## [2,]    1  18   6   75
## [3,]    0  10   0   50
```

```
# If you want to find the number of rows and columns of the matrix:
nrow(mat)
```

```
## [1] 3
```

```
ncol(mat)
```

```
## [1] 4
```

```
# The dim() function finds both and outputs them as a vector,
# starting with the rows (4) and ending with the columns (3):
dim(mat)
```

```
## [1] 3 4
```

```
# Remember, the output of ANY function can be assigned to an object:
(x <- dim(mat))
```

```
## [1] 3 4
```

```
# Note that, right now, your matrix has no row and column names:
mat
```

```
##      [,1] [,2] [,3] [,4]
## [1,]    0  15   2  126
## [2,]    1  18   6   75
## [3,]    0  10   0   50
```

```
rownames(mat)
```

```
## NULL
```

```
colnames(mat)
```

```
## NULL
```

```
# You can assign vectors of names to these rows and columns:
```

```
rownames(mat) <- c("Benjamin", "Kathryn", "Jonathan")
```

```
colnames(mat) <- c("sex", "age", "delinquencies", "ses")
```

```
# Now, when you print the matrix, it will show the names:
```

```
mat
```

```
##           sex age delinquencies ses
## Benjamin   0  15             2 126
## Kathryn    1  18             6  75
## Jonathan   0  10             0  50
```

```
# Using the as.data.frame() function you can convert this matrix
```

```
# into a 'data frame', R's version of a dataset:
```

```
as.data.frame(mat)
```

```
##           sex age delinquencies ses
## Benjamin   0  15             2 126
## Kathryn    1  18             6  75
## Jonathan   0  10             0  50
```

```
# However, I would recommend learning to work with
```

```
# 'tibbles', the tidyverse equivalent of R's data frame:
```

```
as_tibble(mat)
```

```
## # A tibble: 3 x 4
```

```
##       sex   age delinquencies   ses
##   <dbl> <dbl>         <dbl> <dbl>
## 1     0    15             2    126
## 2     1    18             6     75
## 3     0    10             0     50
```

```
# Let's assign this tibble to an object named 'df':
```

```
df <- as_tibble(mat, rownames = "name")
```

```
# Now, we can save these data as a comma-separated values (.csv) file:
```

```
write.csv(df, file = "my_data.csv")
```

```
# or as an excel spreadsheet (.xlsx):
```

```
# install.packages("readxl")
```

```
library(openxlsx)
```

```
## Warning: package 'openxlsx' was built under R version 4.3.3
```



```
write.xlsx(df, file = "my_data.xlsx")
```

```
# or as a Stata dataset (.dta):  
# install.packages("haven")  
library(haven)
```

```
## Warning: package 'haven' was built under R version 4.3.3
```

```
write_dta(df, "my_data.dta")
```

```
# or as a Feather (.feather):  
# install.packages("feather")  
library(feather)
```

```
## Warning: package 'feather' was built under R version 4.3.3
```

```
write_feather(df, "my_data.feather")
```

4 Logical Operations

```
# Logical operations are useful for identifying specific case(s) in your data.  
# So, let's load our data back into R:  
read_csv("my_data.csv")
```

```
## New names:  
## Rows: 3 Columns: 6  
## -- Column specification  
## ----- Delimiter: "," chr  
## (1): name dbl (5): ...1, sex, age, delinquencies, ses  
## i Use 'spec()' to retrieve the full column specification for this data. i  
## Specify the column types or set 'show_col_types = FALSE' to quiet this message.  
## * ' -> '...1'
```

```
## # A tibble: 3 x 6  
##   ...1 name      sex  age delinquencies  ses  
##   <dbl> <chr>    <dbl> <dbl>         <dbl> <dbl>  
## 1     1 Benjamin      0   15             2  126  
## 2     2 Kathryn      1   18             6   75  
## 3     3 Jonathan      0  10             0   50
```

```
read_dta("my_data.dta")
```

```
## # A tibble: 3 x 5  
##   name      sex  age delinquencies  ses  
##   <chr>    <dbl> <dbl>         <dbl> <dbl>  
## 1 Benjamin      0   15             2  126  
## 2 Kathryn      1   18             6   75  
## 3 Jonathan      0  10             0   50
```

```
read_feather("my_data.feather")
```

```
## # A tibble: 3 x 5
##   name      sex  age delinquencies  ses
##   <chr>    <dbl> <dbl>         <dbl> <dbl>
## 1 Benjamin    0   15             2   126
## 2 Kathryn     1   18             6    75
## 3 Jonathan    0   10             0    50
```

What have we forgotten to do here? Assign it to an object!

```
df <- read.xlsx("my_data.xlsx")
```

The openxlsx package will read the data as a data frame, not a tibble.

```
df
```

```
##      name sex age delinquencies ses
## 1 Benjamin  0  15             2 126
## 2 Kathryn  1  18             6  75
## 3 Jonathan  0  10             0  50
```

*# You will need to convert it back into a tibble manually. You can do this
in the same way you did previously, using the as_tibble() function.*

```
df <- as_tibble(df)
```

*# You can combine these functions using the "pipe operator" (%>% or |>).
This operator functions a little like a mathematical operation,
but it is instead used to chain multiple functions.
The chained functions are executed from left to right:*

```
read.xlsx("my_data.xlsx") %>% as_tibble() # tidyverse (magrittr)
```

```
## # A tibble: 3 x 5
##   name      sex  age delinquencies  ses
##   <chr>    <dbl> <dbl>         <dbl> <dbl>
## 1 Benjamin    0   15             2   126
## 2 Kathryn     1   18             6    75
## 3 Jonathan    0   10             0    50
```

```
read.xlsx("my_data.xlsx") |> as_tibble() # base R
```

```
## # A tibble: 3 x 5
##   name      sex  age delinquencies  ses
##   <chr>    <dbl> <dbl>         <dbl> <dbl>
## 1 Benjamin    0   15             2   126
## 2 Kathryn     1   18             6    75
## 3 Jonathan    0   10             0    50
```

*# Now, let's read the .xlsx file, convert it into a tibble, assign
the tibble to an object, and print the tibble. All at the same time!*

```
(df <- read.xlsx("my_data.xlsx") |> as_tibble())
```

```
## # A tibble: 3 x 5
##   name      sex  age delinquencies  ses
##   <chr>   <dbl> <dbl>         <dbl> <dbl>
## 1 Benjamin     0   15             2   126
## 2 Kathryn      1   18             6    75
## 3 Jonathan     0   10             0    50

# ===== #

# Having loaded in our dataset, let's perform some logical operations.
# First, we extract the variable on which we want to perform the operations:
(age <- df$age)

## [1] 15 18 10

# Now, let's check to see if any of the respondents are 18 years old:
age == 18

## [1] FALSE TRUE FALSE

# Are any respondents NOT 18 years old?
age != 18

## [1] TRUE FALSE TRUE

# Are any respondents older than 10 years?
age > 10

## [1] TRUE TRUE FALSE

# Are any respondents less than or equal to 15 years?
age <= 15

## [1] TRUE FALSE TRUE

# Are any respondents less than 12 years OR greater than 17 years?
age < 12 | age > 17

## [1] FALSE TRUE TRUE

# Are any respondents greater than 12 AND less than 17?
age > 12 & age < 17

## [1] TRUE FALSE FALSE

# So far, all of these logical operations have produced
# logical (TRUE / FALSE) vectors. Instead, we might want to
# know the position within the data. The which() function
# can be used to return a vector of rows (indices) that
# fulfil the conditions of the logical operation:
which(age > 12 & age < 17)
```

```
## [1] 1
```

```
# You can invert (negate) these logical operations with a  
# leading exclamation mark (!):  
!(age > 12 & age < 17)
```

```
## [1] FALSE TRUE TRUE
```

```
# This also works if you assign the logical vector:  
(logic <- age > 12 & age < 17)
```

```
## [1] TRUE FALSE FALSE
```

```
!logic
```

```
## [1] FALSE TRUE TRUE
```

```
# The %in% logical operator is unique in that it can be used  
# to compare multiple vectors of length > 1.
```

```
# Is the number '20' in the age vector?  
20 %in% age
```

```
## [1] FALSE
```

```
# What about the number '18'?  
18 %in% age
```

```
## [1] TRUE
```

```
# Are any of the values in the age vector also in  
# a vector consisting of the numbers 10 and 18?  
age %in% c(10, 18)
```

```
## [1] FALSE TRUE TRUE
```

```
# These logical operators can be combined to create  
# increasingly complicated logical statements.  
which(!(age %in% c(10, 18)))
```

```
## [1] 1
```

```
# Just make sure you interpret them carefully and correctly!  
# In this case, we checked to see WHICH respondents' ages  
# do NOT (!) appear in the vector: [10, 18].  
  
# Logical vectors can be interpreted as a binary variable.  
# They are simply encoded as "FALSE" and "TRUE", rather than 0 and 1.  
# As a result, they are easily converted into a binary numeric vector:  
as.numeric(age == 18)
```

```
## [1] 0 1 0
```

```
# So, if you want to count how many respondents fulfill the condition:  
sum(age == 18)
```

```
## [1] 1
```

```
# Or, if you want to find the proportion of respondents who fulfill  
# the condition:  
mean(age == 18)
```

```
## [1] 0.3333333
```

```
# For a percentage, just multiply by 100:  
mean(age == 18) * 100
```

```
## [1] 33.33333
```

```
# For a rate per 1,000, just multiply by 1,000:  
mean(age == 18) * 1000
```

```
## [1] 333.3333
```

```
# ===== #  
  
# For the following exercises and the next section, we are going to load  
# in one of R's "built-in" practice datasets, "USArrests", a data frame  
# reporting violent crime rates in each US state:  
data("USArrests")  
# In this special case, we do not need to assign the data() function, because  
# this function is simply importing the USArrests object from R's files.  
  
# Let's look at the first 6 observations using the head() function:  
head(USArrests)
```

```
##           Murder Assault UrbanPop Rape  
## Alabama      13.2      236      58 21.2  
## Alaska       10.0      263      48 44.5  
## Arizona       8.1      294      80 31.0  
## Arkansas      8.8      190      50 19.5  
## California    9.0      276      91 40.6  
## Colorado     7.9      204      78 38.7
```

```
# Let's look at the first 20 observations instead:  
head(USArrests, 20)
```

```
##           Murder Assault UrbanPop Rape  
## Alabama      13.2      236      58 21.2  
## Alaska       10.0      263      48 44.5  
## Arizona       8.1      294      80 31.0
```

```
## Arkansas      8.8      190      50 19.5
## California    9.0      276      91 40.6
## Colorado      7.9      204      78 38.7
## Connecticut   3.3      110      77 11.1
## Delaware      5.9      238      72 15.8
## Florida       15.4     335      80 31.9
## Georgia       17.4     211      60 25.8
## Hawaii        5.3       46      83 20.2
## Idaho         2.6      120      54 14.2
## Illinois      10.4     249      83 24.0
## Indiana       7.2      113      65 21.0
## Iowa         2.2       56      57 11.3
## Kansas        6.0      115      66 18.0
## Kentucky      9.7      109      52 16.3
## Louisiana     15.4     249      66 22.2
## Maine         2.1       83      51 7.8
## Maryland     11.3     300      67 27.8
```

```
# "USArrests" is a little long, so let's reassign it to the "df" object
df <- USArrests

# Now we check to see if the data were assigned:
head(df)
```

```
##           Murder Assault UrbanPop Rape
## Alabama      13.2     236      58 21.2
## Alaska       10.0     263      48 44.5
## Arizona       8.1     294      80 31.0
## Arkansas      8.8     190      50 19.5
## California    9.0     276      91 40.6
## Colorado      7.9     204      78 38.7
```

5 Performing Simple Statistical Functions

```
# Say we're interested in calculating descriptive statistics for murder.
# First, let's extract the murder variable from the dataset:
murder <- df$Murder
murder
```

```
## [1] 13.2 10.0 8.1 8.8 9.0 7.9 3.3 5.9 15.4 17.4 5.3 2.6 10.4 7.2 2.2
## [16] 6.0 9.7 15.4 2.1 11.3 4.4 12.1 2.7 16.1 9.0 6.0 4.3 12.2 2.1 7.4
## [31] 11.4 11.1 13.0 0.8 7.3 6.6 4.9 6.3 3.4 14.4 3.8 13.2 12.7 3.2 2.2
## [46] 8.5 4.0 5.7 2.6 6.8
```

```
# Median
median(murder)
```

```
## [1] 7.25
```

```
# Mean  
mean(murder)
```

```
## [1] 7.788
```

```
# Standard Deviation  
sd(murder)
```

```
## [1] 4.35551
```

```
# Minimum and Maximum  
min(murder)
```

```
## [1] 0.8
```

```
max(murder)
```

```
## [1] 17.4
```

```
# Frequency Tables  
## Let's create a nominal variable to work with:  
sex <- c(rep("male", 15), rep("female", 8))  
sex
```

```
## [1] "male" "male" "male" "male" "male" "male" "male" "male"  
## [9] "male" "male" "male" "male" "male" "male" "male" "female"  
## [17] "female" "female" "female" "female" "female" "female" "female"
```

```
## Now let's tabulate the results:  
table(sex)
```

```
## sex  
## female male  
##      8   15
```

```
# Missing Data  
# R does not always "know" what to do with missing data.  
# Let's introduce some missingness to the murder variable:  
murder <- c(murder, rep(NA, 5))  
murder
```

```
## [1] 13.2 10.0 8.1 8.8 9.0 7.9 3.3 5.9 15.4 17.4 5.3 2.6 10.4 7.2 2.2  
## [16] 6.0 9.7 15.4 2.1 11.3 4.4 12.1 2.7 16.1 9.0 6.0 4.3 12.2 2.1 7.4  
## [31] 11.4 11.1 13.0 0.8 7.3 6.6 4.9 6.3 3.4 14.4 3.8 13.2 12.7 3.2 2.2  
## [46] 8.5 4.0 5.7 2.6 6.8 NA NA NA NA NA
```

```
# Now let's try to calculate the mean again:  
mean(murder)
```

```
## [1] NA
```

```
# It doesn't work! You have to tell R whether or not to remove
# missing data from the calculation. Under the default setting
# for some functions (e.g., mean), R will try to complete
# calculations with the "NA" values included. As a result,
# the mean cannot be computed. To remove "NA" values, simply
# set the na.rm option to "TRUE":
```

```
mean(murder, na.rm = TRUE)
```

```
## [1] 7.788
```

```
# Similarly, R does not know how to perform logical functions
# on missing values:
```

```
murder > 5
```

```
## [1] TRUE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE TRUE TRUE FALSE
## [13] TRUE TRUE FALSE TRUE TRUE TRUE FALSE TRUE FALSE TRUE FALSE TRUE
## [25] TRUE TRUE FALSE TRUE FALSE TRUE TRUE TRUE TRUE TRUE FALSE TRUE TRUE
## [37] FALSE TRUE FALSE TRUE FALSE TRUE TRUE FALSE FALSE TRUE FALSE TRUE
## [49] FALSE TRUE NA NA NA NA NA
```

```
# You can check for NA's using the is.na() logical function:
```

```
is.na(murder)
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [25] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [37] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [49] FALSE FALSE TRUE TRUE TRUE TRUE TRUE
```

```
# Counting missingness in the variable:
```

```
sum(is.na(murder))
```

```
## [1] 5
```

```
murder |> is.na() |> sum()
```

```
## [1] 5
```

```
murder %>% is.na %>% sum
```

```
## [1] 5
```

```
# Proportion of missingness in the variable:
```

```
mean(is.na(murder))
```

```
## [1] 0.09090909
```



```
murder |> is.na() |> mean()
```

```
## [1] 0.09090909
```

```
murder %>% is.na %>% mean
```

```
## [1] 0.09090909
```

```
# Percent missingness:  
mean(is.na(murder)) * 100
```

```
## [1] 9.090909
```

```
(murder |> is.na() |> mean()) * 100
```

```
## [1] 9.090909
```

```
(murder %>% is.na %>% mean) * 100
```

```
## [1] 9.090909
```

6 Creating Custom Functions

```
# R is a language that is based on functions. mean(), sd(), min(), and every  
# other task you have asked R to perform is based on an underlying function  
# stored within the base R installation, or an installed R package.
```

```
# You can print the underlying function by entering the function call  
# without the succeeding parentheses.
```

```
# So, to print the 'mean' function, you would simply enter:  
mean
```

```
## function (x, ...)  
## UseMethod("mean")  
## <bytecode: 0x000002d762aab4f8>  
## <environment: namespace:base>
```

```
# If this looks like nonsense to you, don't worry, it should.  
# The base R mean function does not use R syntax.  
# So, let's build a function that DOES use R syntax!
```

```
mymean <- function(input, rm_na = FALSE) {  
  x <- sum(input, na.rm = rm_na)  
  if (rm_na == TRUE){  
    y <- sum(!is.na(input))  
  } else {
```

```
y <- length(input)
}
result <- x / y
return(result)
}

# Now let's test the function!
mymean(murder)
```

```
## [1] NA
```

```
# Oops! Forgot to remove missing values with the "rm_na" option
# we included in the custom function!
mymean(murder, rm_na = TRUE)
```

```
## [1] 7.788
```

```
# Compare to the native function:
mean(murder, na.rm = TRUE)
```

```
## [1] 7.788
```

```
# Then, if we want to print our function in the console:
mymean
```

```
## <srcref: file "" chars 14:11 to 23:1>
## <bytecode: 0x000002d76c4e6338>
```